



SURF

Embarrassingly parallel workflows

Autumn HPC school, TU/e

17.11.2025

Embarrassingly Parallel workflows and how to execute them

Embarrassingly parallel workflows

What?



- Executing many independent small jobs simultaneously to use the hardware resource efficiently.

When?

- It is not possible to combine multiple runs/tasks into a single application.
- Large parameter sweep based on the same application or workflow with dependent pipeline steps.
- Workflow involves runs that are depending on several independent software which do not scale beyond a certain number of cores.

Why?

- Nodes are getting denser with respect to resources such as number of cores/threads, GPU cards, Memory.
- Applications do not necessarily scale even within a node: shared memory (OpenMP, MPI SHM, pthreads).



How?

Tools

- Linux background processes (within a node) along with **numactl**.
 - Supported on Snellius and requires prior knowledge of bash scripting and process placement using **numactl**.
 - Provides maximum manual control if one knows what to do.
- Slurm arrays (to perform similar jobs on multiple allocations).
 - Supported on Snellius and requires prior knowledge of SLURM array options.
 - Has some limitations.
- Services for execution of many jobs inside a single allocation, e.g. QCG-PilotJob (<https://qcg-pilotjob.readthedocs.io/en/develop/>).
 - Supported on Snellius and requires prior knowledge of Python scripting.



How?

Tools

- Linux background processes (within a node) along with **numactl**.
 - Supported on Snellius and requires prior knowledge of bash scripting and process placement using **numactl**.
 - Provides maximum manual control if one knows what to do.
- Slurm arrays (to perform similar jobs on multiple allocations).
 - Supported on Snellius and requires prior knowledge of SLURM array options.
 - Has some limitations.
- Services for execution of many jobs inside a single allocation, e.g. QCG-PilotJob (<https://qcg-pilotjob.readthedocs.io/en/develop/>).
 - Supported on Snellius and requires prior knowledge of Python scripting.



| There are many more tools



HyperQueue



PyCOMPSs



PiCaS pilot framework

Slurm Arrays

- Slurm job arrays offer a mechanism for submitting and managing **collections of similar jobs**.
- All jobs must have the **same initial options** (e.g. size, time limit, etc.).
- The **max array size** depends on the settings defined in the SLURM configuration file. On Snellius it's **30000**.
- The **allocation size** per instance **cannot be below the minimum allocation** within SLURM. For instance, on Snellius it's 16 cores and 28 GiB of memory on the “rome” partition.
- One must **wait for each task** in the array to get allocated and finish.


```
#!/bin/bash  
#SBATCH -N 1  
#SBATCH -n 2  
#SBATCH -p rome  
#SBATCH -t 1  
hostname
```

job.sh

```
$ sbatch --array=0-10 job.sh
```



Example: Inefficient use of Slurm Arrays




Linux background process

```
for block in seq 1 10  
do  
  for i in seq 1 16 do  
    serial_program &  
  done  
wait  
done
```


Send tasks to the
background

Wait for all tasks in the
inner loop



Managed workflows

```
...  
from qcg.pilotjob.api.job import Jobs  
  
jobs = Jobs()  
jobs.add(...)  
jobs_ids = manager.submit(jobs)  
print('submitted jobs: ', str(job_ids))  
...
```



| Slurm arrays

Ways of control

- Jobs that are part of a job array will have the environment variable **SLURM_ARRAY_TASK_ID** set to its array index value.
- A **maximum number of simultaneously running tasks** from the job array may be specified using a **"%" separator**.

```
$ cat job.sh
#!/bin/bash
#SBATCH -N 1
#SBATCH -n 2
#SBATCH -p rome
#SBATCH -t 1
echo "SLURM_ARRAY_TASK_ID: $SLURM_ARRAY_TASK_ID"
hostname

$ sbatch --array=0-10 job.sh
...
$ sbatch --array=0-10%2 job.sh
...
```



Ways of control

- Jobs that are part of a job array will have the environment variable **SLURM_ARRAY_TASK_ID** set to its array index value
- A **maximum number of simultaneously running tasks** from the job array may be specified using a **"%" separator**.

```
$ cat job.sh
#!/bin/bash
#SBATCH -N 1
#SBATCH -n 2
#SBATCH -p rome
#SBATCH -t 1
echo "SLURM_ARRAY_TASK_ID: $SLURM_ARRAY_TASK_ID"
hostname
```

```
$ sbatch --array=0-10 job.sh
...
$ sbatch --array=0-10%2 job.sh
...
```

Returns a task ID

Will tell SLURM to submit at most 2 tasks at the same time



Ways of control

Useful environment variables

- **SLURM_ARRAY_JOB_ID** will be set to the first job ID of the array.
- **SLURM_ARRAY_TASK_ID** will be set to the job array index value.
- **SLURM_ARRAY_TASK_COUNT** will be set to the number of tasks in the job array.
- **SLURM_ARRAY_TASK_MAX** will be set to the highest job array index value.
- **SLURM_ARRAY_TASK_MIN** will be set to the lowest job array index value.



Ways of control

File names

Two additional options are available to specify a job's **stdin**, **stdout**, and **stderr** file names:

- **%A** will be replaced by the value of **SLURM_ARRAY_JOB_ID**
- **%a** will be replaced by the value of **SLURM_ARRAY_TASK_ID**

```
#!/bin/bash
...
#SBATCH --output slurm-%A_%a.out
...
```



| Hands-on

Hands-on

Clone the repository

- <https://github.com/sara-nl/course-module-qcg-pilotjob>



```
$ cd ~  
$ git clone https://github.com/sara-nl/course-module-qcg-pilotjob  
$ cd course-module-qcg-pilotjob
```

Hands-on #1

Task

- Submit a job array and print **hostname**, **TMPDIR**, **SLURM_ARRAY_JOB_ID**

Steps

- Go to **hands-on/example_01**
- Take a look at the job script **array_job.sh**
- Submit the SLURM job script, check the SLURM log file

Questions

- How many output files were created?
- Are host names the same?
- What are the names of **TMPDIR** that are reported?

```
#!/bin/bash
#SBATCH -N 1
#SBATCH -n 1
#SBATCH -p rome
#SBATCH -t 10
#SBATCH --array 0-11%3
```

```
echo "Running on "
hostname
echo "TMPDIR $TMPDIR"
```

```
echo "SLURM_ARRAY_JOB_ID $SLURM_ARRAY_JOB_ID"
echo "SLURM_ARRAY_TASK_ID $SLURM_ARRAY_TASK_ID"
echo "SLURM_ARRAY_TASK_COUNT $SLURM_ARRAY_TASK_COUNT"
echo "SLURM_ARRAY_TASK_MAX $SLURM_ARRAY_TASK_MAX"
echo "SLURM_ARRAY_TASK_MIN $SLURM_ARRAY_TASK_MIN"
```

Hands-on #2

Task

- Compute the average of one column in 1000 CSV files
- Aggregate the results in one CSV file.

Steps

- Go to **hands-on/example_02**
- Take a look at the job script **array_job.sh**
- Run **prepare_input.sh**
- Submit the SLURM job script
- Run **aggregate.sh** when all jobs are finished

Questions

- What is the total execution time and SBUs spent on this task?

```
#!/bin/bash
#SBATCH -N 1
#SBATCH -n 1
#SBATCH -p rome
#SBATCH -t 10
#SBATCH --array 0-999%10
```

```
module load 2024
module load Python/3.12.3-GCCcore-13.3.0
```

```
# copy input data and scripts to working directory
cp -r $SLURM_SUBMIT_DIR/input_mod $TMPDIR/
cp -r $SLURM_SUBMIT_DIR/average.py $TMPDIR/
```

```
cd $TMPDIR
```

```
echo "TMPDIR = $TMPDIR"
```

```
# average data
input_file_path=input_mod/${SLURM_ARRAY_TASK_ID}_*
echo "Analyzing file: $input_file_path"
python3 average.py $input_file_path
```

| QCG-PilotJob

QCGPilot-Job


Job manager inside SLURM

- Submits jobs within a single large allocation across multiple nodes.
- Does not have to wait for small individual allocations like in SLURM arrays.
- Tries to optimize resources for each job based on requirements such that the queue is always full.
- Possibility of restarting failed/timed out jobs.
- One can also specify complex result-based job dependencies.

Some caveats and drawbacks

- One needs to estimate, the total time of all the jobs combined.
- Job/system state files get produced in the submission directory.
- Might bring some complexity into execution of a Singularity containers with embedded MPI-based codes.
- In the case of problems, the log files may not be very expressive.

Workflow




QCGPilot-Job: Set up the environment

- The **QCG-PilotJob** package is available as a module in the **2024** environment.
- Note that the environment from the login node is ***not transferred*** to the compute nodes, unless defined in, e.g., the **.bashrc** file. Thus, you have to re-load the QCG-PilotJob package in your job script.

```
$ module purge  
$ module load 2024  
$ module load QCG-PilotJob/0.14.1-foss-2024a
```

Workflow



QCGPilot-Job: SLURM job script

run.sh

```
#!/bin/bash
```

```
#SBATCH -N 1
```

```
#SBATCH -n 16
```

```
#SBATCH -p rome
```


```
#SBATCH -t 10
```

```
module load 2024
```

```
module load QCG-PilotJob/0.14.1-gfbf-2024a
```

```
python test.py
```

Workflow



QCGPilot-Job: Python script

- Work with the **QCG-PilotJob** manager consists of a few mandatory steps:
 - Creating an instance of the **QCG-PilotJob Manager**
 - creation of a class object
 - Waiting for all submitted jobs to finish
 - will wait until some or all submitted to the **QCG-PilotJob Manager** jobs finish
 - Calling for the **finish()** method at the end of the Python script
 - this method terminates the background thread in which the instance of a **QCG-PilotJob Manager** has been run

QCGPilot-Job: Python script

test.py

```
from qcg.pilotjob.api.manager import LocalManager

manager = LocalManager()
print('available resources: ', manager.resources())
manager.wait4all()
manager.finish()
```

Example output

```
available resources: {'total_nodes': 1, 'total_cores': 96, 'used_cores': 0,
'free_cores': 96}
```

QCGPilot-Job: Python script

test.py

Create an object

Wait for all jobs to finish

Finalize the execution

```
from qcg.pilotjob.api.manager import LocalManager

manager = LocalManager()
print('available resources: ', manager.resources())
manager.wait4all()
manager.finish()
```

Example output

```
available resources: {'total_nodes': 1, 'total_cores': 96, 'used_cores': 0,
'free_cores': 96}
```

QCGPilot-Job: Python script

test.py

Create an object

Wait for all jobs to finish

Finalize the execution

```
from qcg.pilotjob.api.manager import LocalManager

manager = LocalManager()
print('available resources: ', manager.resources())
manager.wait4all()
manager.finish()
```

Example output

```
available resources: {'total_nodes': 1, 'total_cores': 96, 'used_cores': 0,
'free_cores': 96}
```

Available cores

Total number of nodes and cores within an allocation

SURF

QCGPilot-Job: Create jobs

- To execute multiple jobs within a single SLURM allocation we need to:
 - Create an object of the **Jobs** class
 - Populate it with some parameters (describe jobs)
 - Submit the created object to the **QCG-PilotJob Manager**

```
...  
from qcg.pilotjob.api.job import Jobs  
  
jobs = Jobs()  
jobs.add(...)  
jobs_ids = manager.submit(jobs)  
print('submitted jobs: ', str(job_ids))  
...
```

QCGPilot-Job: Job description (main parameters)

- **name** - the job name
- **exec** - path to the executable program
- **args** - executable program arguments
- **script** - bash script content
- **stdin** - path to file which content should be passed to the standard input stream
- **stdout** - path to the file where standard output stream should be saved
- **stderr** - path to the file where standard error stream should be saved
- **modules** - list of modules that should be loaded before job start
- **numCores** - number of required cores specification
- **numNodes** - number of required nodes specification
- **iteration** - iterations definition
- **model** - model of execution

QCGPilot-Job: Job description (main parameters)

- **name** - the job name
- **exec** - path to the executable program
- **args** - executable program arguments
- **script** - bash script content
- **stdin** - path to file which content should be passed to the standard input stream
- **stdout** - path to the file where standard output stream should be saved
- **stderr** - path to the file where standard error stream should be saved
- **modules** - list of modules that should be loaded before job start
- **numCores** - number of required cores specification
- **numNodes** - number of required nodes specification
- **iteration** - iterations definition
- **model** - model of execution

QCGPilot-Job: execution models

- **default** - only a single process is launched within the allocation
- **threads** - is designed for running OpenMP tasks on a single node
- **openmpi** - the processes are started with the mpirun command
- **intelmpi** - same but for Intel MPI
- **srunmpi** - the processes are started with the srun command

Read more about execution models in the official docs:

https://qcg-pilotjob.readthedocs.io/en/develop/execution_models.html

QCGPilot-Job: Calling for an executable

test.py

```
from qcg.pilotjob.api.manager import LocalManager
from qcg.pilotjob.api.job import Jobs

manager = LocalManager()
jobs = Jobs()

jobs.add(script='echo "job ${it} executed at `date` @ `hostname`"',
         name='test_job',
         stdout='job.out.${it}',
         iteration=4)

job_names = manager.submit(jobs)
print('submitted job names: ', str(job_names))
job_status = manager.status(job_names)
print('job status: ', job_status)

manager.wait4all()
manager.finish()
```

| Hands-on

SURF

Hands-on #3

Task

- Execute **512 independent short jobs** on two nodes on the “rome” partition. Note that one node on the “rome” partition has 128 cores.

Steps

- Go to **hands-on/example_03**
- Take a look at **qcg_job.py** and **qcg_job.sh**
- Submit the SLURM job script, check the results

Questions

- What host names are reported?
- What additional files and directories are created?

```
from qcg.pilotjob.api.job import Jobs
from qcg.pilotjob.api.manager import LocalManager
```

```
# create the QCG manager and the Jobs object
manager = LocalManager()
jobs = Jobs()
```

```
# create 512 independent jobs
for job in range(512):
    # define a job
    jobs.add(name="job_{}".format(job),
             exec='hostname',
             stdout='output/job{}.out'.format(job),
             stderr='output/job{}.err'.format(job),
             model='default',
             numCores={ "exact" : 1},
             iteration=1)
```

```
print("-- submit jobs --")
manager.submit(jobs)
print("-- wait for all jobs --")
manager.wait4all()
manager.finish()
print("-- finished")
```

Hands-on #4

Task

- Compute the average of one column in 1000 CSV files.
- Aggregate the results in one CSV file.

Steps

- Go to **hands-on/example_04**
- Take a look at **README.md**, **qcg_job_v0.py** and **qcg_job.sh**
- Extract the input data (**input.tar.gz**)
- Submit the SLURM job script, check the results

Steps

- The ***v0** script has a problem. Can you spot it?
- Execute **qcg_job_v1.py**, compare results with the results from **qcg_job_v0.py**

```
job_names = []
```

```
# scan input directory for CVS files
```

```
for filename in glob(os.path.join("input/*.csv"), recursive = False):
```

```
# determine the job name and append it to the list
```

```
    job_names.append(os.path.basename(filename))
```

```
# create the QCG manager and the Jobs object
```

```
manager = LocalManager()
```

```
jobs = Jobs()
```

```
# loop over all files and populate the list of jobs
```

```
for jname in job_names:
```

```
    print("submit {}".format(jname))
```

```
    jobs.add(name=jname,
```

```
            exec='python3',
```

```
            args=["average.py",
```

```
                  os.path.join("input", jname)],
```

```
            stdout='average_{}'.format(jname),
```

```
            stderr='job.{}.err'.format(jname),
```

```
            modules=[
```

```
                "2023", "Python/3.11.3-GCCcore-12.3.0"],
```

```
            iteration=1
```

```
    )
```

```
# add a job for the aggregation of the results
```

```
jobs.add(name="aggregate",
```

```
        script='cat average_*.csv | sort',
```

```
        stdout='result.csv',
```

```
        stderr='aggregate.err')
```



Hands-on #5


Task

- Compute profile for bash pipelines and QCG pilot jobs.
- Generate plots
- Compute efficiency

Steps

- Go to **hands-on/example_05**
- Take a look at **README.md**, **sleep.py**, **pipeline.sh**, **qcg_job.sh**
- Submit the Slurm job scripts, check the results
- Follow README.md to generate profiles





SURF

Have fun! :)